

class11

Vince (PID: A15422556)

2/22/2022

Q1. How many genes are in this dataset? 38694

Q2. How many 'control' cell lines do we have? 4 'control' cell lines

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")

nrow(counts)
```

```
## [1] 38694
```

```
sum(metadata$dex == "control")
```

```
## [1] 4
```

Check to see that columns of countdata and coldata (metadata) match.

```
all(metadata$id == colnames(counts))
```

```
## [1] TRUE
```

Q3. How would you make the above code in either approach more robust? See code below.

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean) See code below.

Extract control and treated counts for comparison

Extract the control counts columns.

```
control.ids <- metadata[metadata$dex == "control", "id"]
control.counts <- counts[,control.ids]

control.mean <- rowMeans(control.counts)
head(control.mean)
```

```
## ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
##           900.75           0.00           520.50           339.75           97.25
## ENSG00000000938
##           0.75
```

```
treated.ids <- metadata[metadata$dex == "treated", "id"]
treated.counts <- counts[,treated.ids]

treated.mean <- rowMeans(treated.counts)
head(treated.mean)


## ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
##          658.00          0.00          546.00          316.50          78.75
## ENSG00000000938
##          0.00
```

Plot comparing treated vs. control.

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following.

Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot? geom_point()

```
plot(treated.mean, control.mean,
     xlab = "log of drug treated counts",
     ylab = "log of control counts")
```

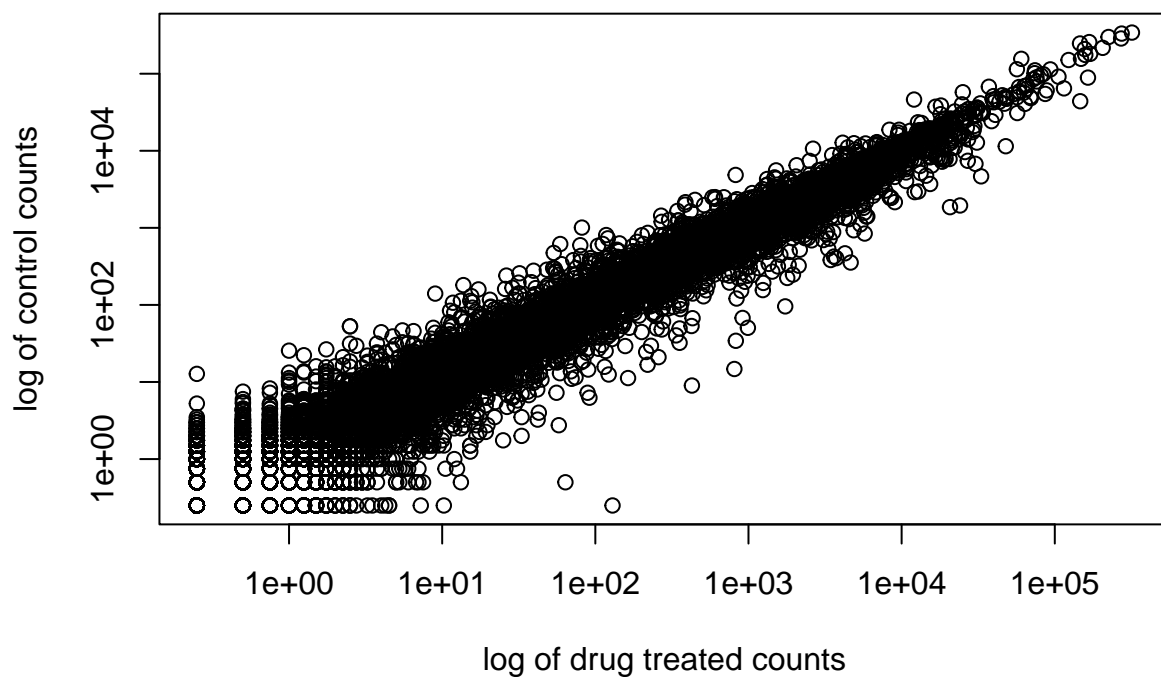


Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this? log="xy"

```
plot(treated.mean, control.mean, log = "xy",
     xlab = "log of drug treated counts",
     ylab = "log of control counts")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 x values <= 0 omitted
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 y values <= 0 omitted
## from logarithmic plot
```



Changes in gene expression: treated vs. control. This would represent points (i.e. genes) that do not lie on the diagonal.

```
log2fc <- log2(treated.mean / control.mean)

meancounts <- data.frame(control.mean, treated.mean, log2fc)
head(meancounts)
```

##	control.mean	treated.mean	log2fc
## ENSG00000000003	900.75	658.00	-0.45303916
## ENSG00000000005	0.00	0.00	NaN
## ENSG000000000419	520.50	546.00	0.06900279
## ENSG000000000457	339.75	316.50	-0.10226805
## ENSG000000000460	97.25	78.75	-0.30441833
## ENSG000000000938	0.75	0.00	-Inf

Q7. What is the purpose of the `arr.ind` argument in the `which()` function call above? Why would we then take the first column of the output and need to call the `unique()` function? Tells the row and column where the values are true.

```
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)

to.rm <- unique(zero.vals[,1])
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

```
##               control.mean treated.mean      log2fc
## ENSG000000000003         900.75        658.00 -0.45303916
## ENSG000000000419         520.50        546.00  0.06900279
## ENSG000000000457         339.75        316.50 -0.10226805
## ENSG000000000460          97.25         78.75 -0.30441833
## ENSG000000000971        5219.00       6687.50  0.35769358
## ENSG00000001036        2327.00       1785.75 -0.38194109
```

```
nrow(mycounts)
```

```
## [1] 21817
```

Q8. Using the `up.ind` vector above can you determine how many up regulated genes we have at the greater than 2 fc level? 250

Q9. Using the `down.ind` vector above can you determine how many down regulated genes we have at the greater than 2 fc level? 21503

Q10. Do you trust these results? Why or why not? No, we need a p-value

“Up” genes

```
sum(mycounts$log2fc > 2, na.rm = TRUE)
```

```
## [1] 250
```

“Down” genes

```
sum(mycounts$log2fc < -2, na.rm = TRUE)
```

```
## [1] 21503
```

Missing the stats (are differences significant):

DESeq2 Analysis

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```
## Loading required package: BiocGenerics
```

```
##
```

```
## Attaching package: 'BiocGenerics'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      IQR, mad, sd, var, xtabs
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      anyDuplicated, append, as.data.frame, basename, cbind, colnames,
```

```
##      dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,
```

```
##      grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,
```

```
##      order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,
```

```
##      rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,
```

```
##      union, unique, unsplit, which.max, which.min
```

```
##
```

```
## Attaching package: 'S4Vectors'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      expand.grid, I, unname
```

```
## Loading required package: IRanges
```

```
##
```

```
## Attaching package: 'IRanges'
```

```
## The following object is masked from 'package:grDevices':
```

```
##
```

```
##      windows
```

```
## Loading required package: GenomicRanges
```

```
## Loading required package: GenomeInfoDb
```

```
## Loading required package: SummarizedExperiment
```

```
## Loading required package: MatrixGenerics
```

```
## Loading required package: matrixStats
```

```
##
## Attaching package: 'MatrixGenerics'

## The following objects are masked from 'package:matrixStats':
##
##   colAlls, colAnyNAs, colAnys, colAvgPerRowSet, colCollapse,
##   colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
##   colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
##   colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
##   colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
##   colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
##   colWeightedMeans, colWeightedMedians, colWeightedSds,
##   colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgPerColSet,
##   rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
##   rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
##   rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
##   rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
##   rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
##   rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
##   rowWeightedSds, rowWeightedVars
```

```
## Loading required package: Biobase
```

```
## Welcome to Bioconductor
##
##   Vignettes contain introductory material; view with
##   'browseVignettes()'. To cite Bioconductor, see
##   'citation("Biobase")', and for packages 'citation("pkgname")'.
```

```
##
## Attaching package: 'Biobase'
```

```
## The following object is masked from 'package:MatrixGenerics':
##
##   rowMedians
```

```
## The following objects are masked from 'package:matrixStats':
##
##   anyMissing, rowMedians
```

Package wants input in a specific way:

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData=metadata,
                              design=~dex)
```

```
## converting counts to integer mode
```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
## design formula are characters, converting to factors
```

```
dds
```

```
## class: DESeqDataSet
## dim: 38694 8
## metadata(1): version
## assays(1): counts
## rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
## ENSG00000283123
## rowData names(0):
## colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
## colData names(4): id dex celltype geo_id
```

Run the DESeq analysis.

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

Results


```
res <- results(dds)
res
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 38694 rows and 6 columns
##      baseMean log2FoldChange      lfcSE      stat      pvalue
##      <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003  747.1942    -0.3507030  0.168246 -2.084470  0.0371175
## ENSG000000000005    0.0000         NA         NA         NA         NA
## ENSG000000000419  520.1342     0.2061078  0.101059  2.039475  0.0414026
## ENSG000000000457  322.6648     0.0245269  0.145145  0.168982  0.8658106
## ENSG000000000460   87.6826    -0.1471420  0.257007 -0.572521  0.5669691
## ...           ...           ...           ...           ...
## ENSG00000283115    0.000000         NA         NA         NA         NA
## ENSG00000283116    0.000000         NA         NA         NA         NA
## ENSG00000283119    0.000000         NA         NA         NA         NA
## ENSG00000283120   0.974916    -0.668258    1.69456 -0.394354  0.693319
## ENSG00000283123    0.000000         NA         NA         NA         NA
##                padj
```

```
##                               <numeric>
## ENSG000000000003  0.163035
## ENSG000000000005      NA
## ENSG000000000419  0.176032
## ENSG000000000457  0.961694
## ENSG000000000460  0.815849
## ...
## ENSG00000283115      NA
## ENSG00000283116      NA
## ENSG00000283119      NA
## ENSG00000283120      NA
## ENSG00000283123      NA
```

Volcano Plot

```
plot( res$log2FoldChange, -log(res$padj),
      xlab="Log2(FoldChange)",
      ylab="-Log(P-value)")
abline(h=-log(0.05), col="gray")
abline(v=c(-2,2), col="gray")
```




Add color to the plots


```
mycols <- rep("gray", nrow(res))

mycols[res$padj < 0.005] <- "red"
mycols[abs(res$log2FoldChange) < 2] <- "gray"

plot(res$log2FoldChange, -log(res$padj), col=mycols)
```



Adding annotation data

To help interpret our results we need to understand what the differentially expressed genes are. A first step is to get the gene names (i.e. gene SYMBOLs).

```
library("AnnotationDbi")
library("org.Hs.eg.db")
```

```
##
```

What DB identifiers can I look up?

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM"      "ALIAS"        "ENSEMBL"      "ENSEMBLPROT"  "ENSEMBLTRANS"
## [6] "ENTREZID"    "ENZYME"       "EVIDENCE"     "EVIDENCEALL"  "GENENAME"
## [11] "GENETYPE"    "GO"           "GOALL"        "IPI"           "MAP"
## [16] "OMIM"        "ONTOLOGY"     "ONTOLOGYALL"  "PATH"          "PFAM"
## [21] "PMID"        "PROSITE"      "REFSEQ"       "SYMBOL"        "UCSCKG"
## [26] "UNIPROT"
```

Use `mapIds()` function to translate between different IDs.

```
res$symbol <- mapIds(org.Hs.eg.db,
                     keys=row.names(res), # Our genenames
                     keytype="ENSEMBL",   # The format of our genenames
                     column="SYMBOL",     # The new format we want to add
                     multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 7 columns
##           baseMean log2FoldChange    lfcSE      stat    pvalue
##           <numeric>    <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195    -0.3507030  0.168246 -2.084470 0.0371175
## ENSG000000000005  0.000000         NA         NA         NA         NA
## ENSG000000000419 520.134160     0.2061078  0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844     0.0245269  0.145145  0.168982 0.8658106
## ENSG000000000460  87.682625    -0.1471420  0.257007 -0.572521 0.5669691
## ENSG000000000938  0.319167    -1.7322890  3.493601 -0.495846 0.6200029
##           padj      symbol
##           <numeric> <character>
## ENSG000000000003  0.163035      TSPAN6
## ENSG000000000005         NA      TNMD
## ENSG000000000419  0.176032      DPM1
## ENSG000000000457  0.961694      SCYL3
## ENSG000000000460  0.815849      C1orf112
## ENSG000000000938         NA      FGR
```

Q11. Run the `mapIds()` function two more times to add the Entrez ID and UniProt accession and GENENAME as new columns called `res$entrez`, `res$uniprot` and `res$genename`. See code below.

```
res$entrez <- mapIds(org.Hs.eg.db,
                    keys=row.names(res), # Our genenames
                    keytype="ENSEMBL",   # The format of our genenames
                    column="ENTREZID",    # The new format we want to add
                    multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
res$uniprot <- mapIds(org.Hs.eg.db,
  keys=row.names(res), # Our genenames
  keytype="ENSEMBL",   # The format of our genenames
  column="UNIPROT",     # The new format we want to add
  multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
res$genename <- mapIds(org.Hs.eg.db,
  keys=row.names(res), # Our genenames
  keytype="ENSEMBL",   # The format of our genenames
  column="GENENAME",   # The new format we want to add
  multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 10 columns
##
```

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
## ENSG00000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
## ENSG00000000005	0.000000	NA	NA	NA	NA
## ENSG000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
## ENSG000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
## ENSG000000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
## ENSG000000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029

```
##
```

	padj	symbol	entrez	uniprot
	<numeric>	<character>	<character>	<character>
## ENSG00000000003	0.163035	TSPAN6	7105	A0A024RCI0
## ENSG00000000005	NA	TNMD	64102	Q9H2S6
## ENSG000000000419	0.176032	DPM1	8813	O60762
## ENSG000000000457	0.961694	SCYL3	57147	Q8IZE3
## ENSG000000000460	0.815849	C1orf112	55732	A0A024R922
## ENSG000000000938	NA	FGR	2268	P09769

```
##
```

	genename
	<character>
## ENSG00000000003	tetraspanin 6
## ENSG00000000005	tenomodulin
## ENSG000000000419	dolichyl-phosphate m..
## ENSG000000000457	SCY1 like pseudokina..
## ENSG000000000460	chromosome 1 open re..
## ENSG000000000938	FGR proto-oncogene, ..

Pathway analysis with R and Bioconductor

Q12. Can you do the same procedure as above to plot the pathview figures for the top 2 down-regulated pathways? See code below.

```
library(pathview)
```

```
## #####  
## Pathview is an open source software package distributed under GNU General  
## Public License version 3 (GPLv3). Details of GPLv3 is available at  
## http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to  
## formally cite the original Pathview paper (not just mention it) in publications  
## or products. For details, do citation("pathview") within R.  
##  
## The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG  
## license agreement (details at http://www.kegg.jp/kegg/legal.html).  
## #####
```

```
library(gage)
```

```
##
```

```
library(gageData)
```

```
data(kegg.sets.hs)
```

```
# Examine the first 2 pathways in this kegg set for humans  
head(kegg.sets.hs, 2)
```

```
## $'hsa00232 Caffeine metabolism'  
## [1] "10" "1544" "1548" "1549" "1553" "7498" "9"  
##  
## $'hsa00983 Drug metabolism - other enzymes'  
## [1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"  
## [9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"  
## [17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"  
## [25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"  
## [33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"  
## [41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"  
## [49] "8824" "8833" "9" "978"
```

Need a vector of fold-change labeled with the names of our genes in ENTREZ format.

```
foldchanges = res$log2FoldChange  
names(foldchanges) = res$entrez  
head(foldchanges)
```

```
##          7105          64102          8813          57147          55732          2268  
## -0.35070302          NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

Run the GAGE analysis passing in our foldchange vector and KEGG genesets we are interested in.

```
# Get the results  
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

Look at what is contained in this keggres results object (i.e. its attributes).

```
attributes(keggres)
```

```
## $names
## [1] "greater" "less" "stats"
```

```
# Look at the first three down (less) pathways
head(keggres$less, 3)
```

```
##
## hsa05332 Graft-versus-host disease 0.0004250461 -3.473346 0.0004250461
## hsa04940 Type I diabetes mellitus 0.0017820293 -3.002352 0.0017820293
## hsa05310 Asthma 0.0020045888 -3.009050 0.0020045888
##
## q.val set.size exp1
## hsa05332 Graft-versus-host disease 0.09053483 40 0.0004250461
## hsa04940 Type I diabetes mellitus 0.14232581 42 0.0017820293
## hsa05310 Asthma 0.14232581 29 0.0020045888
```

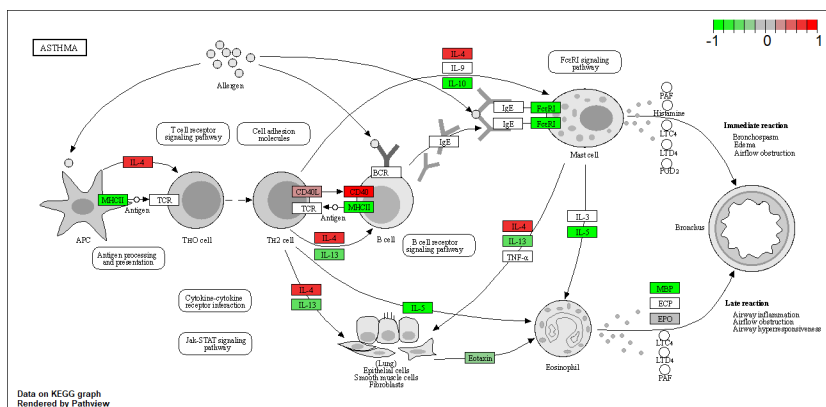
Map my results onto any KEGG pathway. Do this manually first by selecting one of the pathway IDs from above.

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

```
## 'select()' returned 1:1 mapping between keys and columns
```

```
## Info: Working in directory C:/Users/vince/Desktop/UCSD/Academic Years/Fourth Year/BIMM143/class11
```

```
## Info: Writing image file hsa05310.pathview.png
```



Final step is to save our results.

```
write.csv(res, file="deseq_results.csv")
```